

Machine Learning

March 13, 2017

- n sample points $x_i \in \mathbb{R}^d, i = 1, \dots, n$
- $d = 2$ where not stated.


Classification

A **decision boundary** is a curve separating the plane (sample space) into two regions.

Some classifiers involve a **decision function** f , in which case $f(x) = 0$ describes the decision boundary.

A **linear classifier** uses a linear decision function $f(x) = \mathbf{w} \cdot \mathbf{x} + \alpha$. This is scalar-valued: it's a plane over the plane (sample space). Its intersection defines a linear decision boundary.

In d -dimensions the decision boundary is a hyperplane ($(d-1)$ -dimensional). This still separates the sample space into two regions.



Example: $f(x) = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + 4$

- A plane sloping up at 45 in the north-east direction.
- Each input feature has equal influence on the classification.
- Decision boundary is line $x_1 + x_2 = -4$.
- \mathbf{w} is normal to the decision boundary since $\mathbf{w} \cdot (\mathbf{x}_1 - \mathbf{x}_2) = -4 - (-4) = 0$.

If one feature has a very high weight then \mathbf{w} points close to that axis and the decision boundary is almost perpendicular to that axis (other features almost don't matter).

Distance from the decision boundary to a point: For some point \mathbf{x}_i , the height of the decision plane above \mathbf{x}_i is $\mathbf{w} \cdot \mathbf{x}_i + \alpha$. At the decision boundary, this height is zero. Looking "straight up" the slope of the decision function, its gradient is $\sqrt{w_1^2 + w_2^2} = \|\mathbf{w}\|$. So the distance of a point \mathbf{x}_i from the hyperplane is $\frac{\mathbf{w} \cdot \mathbf{x}_i + \alpha}{\|\mathbf{w}\|}$. If \mathbf{w} is not a unit vector, the problem can be rescaled

so that it is, in which case the distance is $\mathbf{w} \cdot \mathbf{x}_i + \alpha$.

Examples of linear classifiers:

- **Centroid method:** Decision boundary perpendicular to and bisects line connecting means of labeled training points.
- **Perceptron:**
- **Maximum margin classifier:**
- **LDA:** Fit Gaussians to each class, same covariance across classes.

Perceptron

Labels $y_i \in \{-1, 1\}$. Assume $\alpha = 0$ for now (decision boundary through origin).

Goal: find line separating points (separating hyperplane). I.e. Find \mathbf{w} such that

$$\begin{cases} \mathbf{x}_i \cdot \mathbf{w} \leq 0, & y_i = -1 \\ \mathbf{x}_i \cdot \mathbf{w} \geq 0, & y_i = +1. \end{cases}$$

This is equivalent to the **constraint** $y_i \mathbf{x}_i \cdot \mathbf{w} \geq 0$.

Cost function: total distance $R(\mathbf{w})$ of misclassified points from the decision boundary.

Optimization problem: Find \mathbf{w} that minimizes

$$R(\mathbf{w}) = \sum_{i \in V} L(\mathbf{x}_i \cdot \mathbf{w}, y_i) = \sum_{i \in V} -y_i \mathbf{x}_i \cdot \mathbf{w},$$

where V are the misclassified points.

Per-training point loss function

$$L(\text{prediction}, y_i) = L(\mathbf{x}_i \cdot \mathbf{w}, y_i) = \begin{cases} 0, & \text{correct, } y_i \mathbf{x}_i \cdot \mathbf{w} \geq 0 \\ -y_i \mathbf{x}_i \cdot \mathbf{w}, & \text{misclassified} \end{cases}$$

Gradient descent: Find \mathbf{w} that minimizes $R(\mathbf{w})$.

$$\nabla_{\mathbf{w}} R = \begin{bmatrix} -\sum_i y_i x_{i1} \\ \vdots \\ -\sum_i y_i x_{id} \end{bmatrix}$$

- On each iteration, compute the gradient; update \mathbf{w} by taking a step downhill of size ρ : $\mathbf{w} \leftarrow \mathbf{w} + \rho \sum_{i \in V} y_i \mathbf{x}_i$.
- A misclassified data point far out in dimension j will cause the gradient to have a large component $-\sum_i y_i x_{ij}$ in that dimension.
- \mathbf{w} thus becomes more closely aligned with that axis and the decision boundary.
- Decision boundary therefore becomes more perpendicular to that axis (axis becomes more "important").


Stochastic gradient descent (Perceptron): on each iteration pick one misclassified point and update \mathbf{w} using gradient for that point: $\mathbf{w} \leftarrow \mathbf{w} + \rho y_i \mathbf{x}_i$.

Allow decision boundaries that do not pass through origin: add a fictitious dimension so that sample points now lie on the plane $x_{d+1} = 1$ in $(d+1)$ dimensions. Run algorithm as above, just with the new dimensionality.

$$\begin{bmatrix} \mathbf{w} \cdot \mathbf{x} + \alpha \\ w_{d+1} \\ w_{d+2} \\ \alpha \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \\ 1 \end{bmatrix} = 0.$$

Optimization in weight space

x-space | **w-space**
hyperplane | point \mathbf{w} is normal vector to hyperplane
point | hyperplane whose normal vector is the \mathbf{x} point (? don't understand this yet)



Maximum margin classifiers

Margin is distance from hyperplane to nearest sample point.

Previously, in the perceptron, we used the constraint

$$y_i \mathbf{x}_i \cdot \mathbf{w} \geq 0.$$

Now, we demand that there is a non-zero margin between the decision boundary and the points:

$$y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1.$$

The 1 on the RHS is arbitrary. I think \mathbf{w} and α will adapt to make it true for any positive value, so the point is that we're demanding a strictly non-zero margin.

Decision Theory ^{3 4}

Suppose there are two possible **classes**: $\{C, D\}$

Decision rule: $r(\mathbf{x}) : \mathbb{R}^d \rightarrow \{C, D\}$

Loss function: E.g. 0-1 loss:

$$L(y_i \rightarrow \hat{y}_i) = \begin{cases} 0, & \hat{y}_i = y_i \\ 1, & \text{otherwise} \end{cases} \quad (\text{correct classification})$$

Risk: Functional $R(r)$: expected loss for rule r , over $p(X, Y)$. ⁵

So what rule function r minimizes the functional R ?

Bayes decision rule: Assign \mathbf{x} to class C if

$$(C \text{ posterior at } \mathbf{x}) \times (\text{penalty for misclassifying a true } C)$$

is largest for class C . I.e. if

$$p(C|\mathbf{x})L(D|C) > p(D|\mathbf{x})L(C|D).$$

With 0-1 loss, this is: "assign to class with highest posterior".

With 0-1 loss and two classes, it's: "assign to class with posterior > 0.5 ".

Empirical risk: Discriminative methods (e.g. logistic regression) lack any model for X . How can we estimate expected loss over $p(X, Y)$? Take the observed sample points as defining a discrete, uniform distribution, in which case

$$\hat{R}(r) = \frac{1}{n} \sum_i L(r(x_i), y_i).$$

This provides a justification for minimizing the sum/mean of per-sample loss.

Statistical justifications

Regression: want to estimate a function f such that $y_i = f(x_i) + \epsilon$, where ϵ has unknown distribution but mean 0. Ideal would be to estimate f with $h(x_i) = \mathbb{E}[Y|x_i]$ since this is equal to $f(x_i)$.

Likelihood justification for linear regression cost function.

Logistic Regression from Maximum Likelihood

Bias-Variance Decomposition

$$\begin{aligned} \text{Bias} &= \mathbb{E}[h(\mathbf{x})] - y^T \\ \text{Variance} &= \mathbb{E}[h(\mathbf{x})^2] - 2\mathbb{E}[h(\mathbf{x})]y^T + y^T y \\ \text{Total Error} &= \mathbb{E}[h(\mathbf{x})^2] - 2\mathbb{E}[h(\mathbf{x})]y^T + y^T y \\ &= \mathbb{E}[h(\mathbf{x}) - y]^2 + \text{Var}(h(\mathbf{x})) + \mathbb{E}[h(\mathbf{x}) - y]^2 - 2\mathbb{E}[h(\mathbf{x})]y^T + y^T y \\ &= \mathbb{E}[h(\mathbf{x}) - y]^2 + \text{Var}(h(\mathbf{x})) + \mathbb{E}[h(\mathbf{x}) - y]^2 - 2\mathbb{E}[h(\mathbf{x})]y^T + y^T y \\ &= \mathbb{E}[h(\mathbf{x}) - y]^2 + \text{Var}(h(\mathbf{x})) + \mathbb{E}[h(\mathbf{x}) - y]^2 - 2\mathbb{E}[h(\mathbf{x})]y^T + y^T y \end{aligned}$$

bias of method variance of method irreducible error

Gaussian discriminant analysis ^{6 7}

Anisotropic:

$$p(\mathbf{x}) = \frac{1}{(2\pi)^{d/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu)^T \Sigma^{-1}(\mathbf{x} - \mu)\right)$$

Isotropic:

$$p(\mathbf{x}) = \frac{1}{(2\pi)^{d/2}\sigma^d} \exp\left(-\frac{\|\mathbf{x} - \mu\|^2}{2\sigma^2}\right)$$

Isotropic Gaussians

Multivariate data \mathbf{x} but features uncorrelated and all features same variance. L

QDA

Fit separate Gaussians to the training data in each class. The likelihood is

$$p(\mathbf{x}|\text{class } C) = \frac{1}{(2\pi)^{d/2}\sigma_C^d} \exp\left(-\frac{\|\mathbf{x} - \mu_C\|^2}{2\sigma_C^2}\right)$$

and we compare the value of $p(\mathbf{x}|\text{class } C) \cdot \pi_C \cdot L(D|C)$.

The decision boundaries are where the posterior \times loss are equal. It's easier to compare the log of this:

$$Q_C(\mathbf{x}) = -\frac{\|\mathbf{x} - \mu_C\|^2}{2\sigma_C^2} - d \log \sigma_C + \log \pi_C + \log L(D|C)$$

The posterior probability of class C at point \mathbf{x} is ⁸

$$p(\mathbf{x}|\mathbf{x}) = \frac{\pi_C p(\mathbf{x}|\mathbf{x})}{\pi_C p(\mathbf{x}|\mathbf{x}) + \pi_D p(\mathbf{x}|\mathbf{x})} = \frac{1}{1 + e^{-(Q_C(\mathbf{x}) - Q_D(\mathbf{x}))}}$$

so logistic in the quadratic expression $Q_C(\mathbf{x}) - Q_D(\mathbf{x})$.

LDA

Estimate separate class means but same variance for all classes. So now

$$\begin{aligned} Q_C(\mathbf{x}) - Q_D(\mathbf{x}) &= \frac{\|\mathbf{x} - \mu_D\|^2 - \|\mathbf{x} - \mu_C\|^2}{\sigma^2} + \log \frac{\pi_C}{\pi_D} + \log \frac{L(D|C)}{L(D|D)} \\ &= \frac{(\mathbf{x} - \mu_D) \cdot (\mathbf{x} - \mu_D) - (\mathbf{x} - \mu_C) \cdot (\mathbf{x} - \mu_C)}{\sigma^2} + \log \frac{\pi_C}{\pi_D} + \log \frac{L(D|C)}{L(D|D)} \\ &= \mathbf{x} \cdot \frac{2(\mu_C - \mu_D)}{\sigma^2} + \left(\frac{\mu_D^2 - \mu_C^2}{\sigma^2} + \log \frac{\pi_C}{\pi_D} + \log \frac{L(D|C)}{L(D|D)} \right) \\ &= \mathbf{x} \cdot \mathbf{w} + \alpha \end{aligned}$$

Optimization problem (quadratic program):
Find \mathbf{w}, α that minimize $\|\mathbf{w}\|^2$ such that $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1$ for all points i .

Soft margin SVMs ^{1 2}

- Still quadratic program but allow points to violate margin via **slack variables** $\xi_i \geq 0$:
- Constraint is $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$.
- Find non-linear decision boundaries by introducing new features comprising non-linear functions of base features ("lift points into higher-dimensional space").

Optimization problem: Find \mathbf{w}, α , and ξ that minimize $\ \mathbf{w}\ ^2 + C \sum_i \xi_i$, subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$ $\xi_i \geq 0$ for all $i = 1, \dots, n$ It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities. C is a scale coefficient that balances the trade-off.	
minimize	$\ \mathbf{w}\ ^2 + C \sum_i \xi_i$
subject to	$y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$
variables	$\mathbf{w}, \alpha, \xi_i$
constraints	$y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$
boundary	min "bar"

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

C is a scale coefficient that balances the trade-off.

minimize $\|\mathbf{w}\|^2 + C \sum_i \xi_i$

subject to: $y_i (\mathbf{x}_i \cdot \mathbf{w} + \alpha) \geq 1 - \xi_i$ for all $i = 1, \dots, n$

$\xi_i \geq 0$ for all $i = 1, \dots, n$

It's a quadratic program because its objective function is quadratic and its constraints are linear inequalities.

³<https://people.eecs.berkeley.edu/~jrs/189/lec/06.pdf>
⁴<https://www.youtube.com/watch?v=KXa8n02iqI8>

⁵ $R(r) = \pi(Y = -1) \mathbb{E}_X L(-1 \rightarrow r(X)) + \pi(Y = +1) \mathbb{E}_X L(+1 \rightarrow r(X))$ over $p(Y)$ $p(X|Y)$

$= \sum_i p(X) (\pi(Y = -1) L(-1 \rightarrow r(X)) + \pi(Y = +1) L(+1 \rightarrow r(X)))$ over $p(X) p(Y|X)$

⁶<https://people.eecs.berkeley.edu/~jrs/189/lec/07.pdf>

⁷<https://www.youtube.com/watch?v=5d6b0e2d2a>

⁸This is assuming 0-1 loss, so the loss doesn't affect $Q_C(\mathbf{x})$

By the definition of eigenvector we have

$$\mathbf{A}\mathbf{V} = \mathbf{V}\mathbf{\Lambda}$$

and therefore the **eigendecomposition** of \mathbf{A}

$$\mathbf{A} = \mathbf{V}\mathbf{\Lambda}\mathbf{V}^T.$$

So we can perform $\mathbf{A}\mathbf{x}$ as $\mathbf{V}\mathbf{\Lambda}\mathbf{V}^T\mathbf{x}$, and $\mathbf{A}^k\mathbf{x}$ as $\mathbf{V}\mathbf{\Lambda}^k\mathbf{V}^T\mathbf{x}$:

1. $\mathbf{V}^T = \mathbf{V}^{-1}$ rotates the input vector into axis-aligned coordinates.
2. $\mathbf{\Lambda}$ scales along different axes.
3. \mathbf{V} returns to the original coordinates.

$\mathbf{\Lambda}$ is said to be the diagonalized version of \mathbf{A} .

9 The Anisotropic Multivariate Normal Distribution, QDA, and LDA

Regression

Linear Least Squares Regression

Use fictitious dimension trick, so that \mathbf{w} includes the offset term α and \mathbf{X} is $(n \times (d + 1))$.

Find \mathbf{w} that minimizes cost function $J(\mathbf{w})$: sum of squared difference between linear predictor and observed training point.

$$J(\mathbf{w}) = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 = \sum_i (\mathbf{x}_i^T \mathbf{w} - y_i)^2$$

Solve by differentiating and finding the critical point:

$$\begin{aligned} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 &= \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y} \\ \nabla_{\mathbf{w}} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 &= 2\mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{X}^T \mathbf{y} \\ \mathbf{w}^* &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} =: \mathbf{X}^+ \mathbf{y} \end{aligned}$$

For a new sample point \mathbf{x} , the prediction is $\hat{y} = \mathbf{x}^+ \cdot \mathbf{w}^*$.

Related concepts

- **normal equations**: linear system of d equations in unknown \mathbf{w} resulting from setting the gradient equal to zero: $\mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y} = 0$
- **pseudoinverse**: The matrix $\mathbf{X}^+ = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$ maps \mathbf{y} to \mathbf{w}^* . In general there's no \mathbf{w} that solves $\mathbf{X} \mathbf{w} = \mathbf{y}$, but $\mathbf{w}^* = \mathbf{X}^+ \mathbf{y}$ makes the LHS as close as possible to \mathbf{y} . So it behaves as a “left inverse” of \mathbf{X} , since $\mathbf{X}^+ \mathbf{X} = \mathbf{I}$ and left-multiplying by \mathbf{X}^+ gives the “solution” to $\mathbf{X} \mathbf{w} = \mathbf{y}$.
- **projection matrix or hat matrix**: Still focusing on the training phase, the predictions are $\hat{\mathbf{y}} = \mathbf{X} \mathbf{w}^* = \mathbf{X} \mathbf{X}^+ \mathbf{y}$. So $\mathbf{X} \mathbf{X}^+$ “puts that hat on \mathbf{y} , or projects \mathbf{y} onto the hyperplane, in the viewpoint described below.

Projection interpretation

Usually we think of n points in \mathbb{R}^d . But instead, consider a separate column of the data for each feature: these are d points in \mathbb{R}^n . The observed training data \mathbf{y} is also a point in \mathbb{R}^n , and so is the prediction $\hat{\mathbf{y}} = \mathbf{X} \mathbf{w}^*$.

As we vary \mathbf{w} , the prediction $\mathbf{X} \mathbf{w}$ describes a hyperplane spanned by the columns of \mathbf{X} .

We want to find the \mathbf{w}^* corresponding to the closest point on the hyperplane to \mathbf{y} . So $\mathbf{X} \mathbf{w}^* - \mathbf{y}$ must be orthogonal to the hyperplane:

$$\mathbf{X}^T \cdot (\mathbf{X} \mathbf{w}^* - \mathbf{y}) = 0.$$

Which are the normal equations (linear system of d equations), derived differently.

Weighted linear regression

Sample point i has weight b_i . Diagonal $n \times n$ matrix \mathbf{B} contains weights.

$$\begin{aligned} J(\mathbf{w}) &= \sum_i b_i (\mathbf{x}_i^T \mathbf{w} - y_i)^2 \\ &= (\mathbf{X} \mathbf{w} - \mathbf{y})^T \mathbf{B} (\mathbf{X} \mathbf{w} - \mathbf{y}) \\ &= \mathbf{w}^T \mathbf{X}^T \mathbf{B} \mathbf{X} \mathbf{w} - 2\mathbf{y}^T \mathbf{B} \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y} \end{aligned}$$

Gradient

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = 2\mathbf{X}^T \mathbf{B} \mathbf{X} \mathbf{w} - 2\mathbf{X}^T \mathbf{B} \mathbf{y}$$

Solution

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{B} \mathbf{X})^{-1} \mathbf{X}^T \mathbf{B} \mathbf{y}$$

How to compute the gradient

The cost function is $J(\mathbf{w}) = \|\mathbf{X} \mathbf{w} - \mathbf{y}\|^2$. We could write this as a dot product and multiply out:

$$\begin{aligned} J(\mathbf{w}) &= (\mathbf{X} \mathbf{w} - \mathbf{y}) \cdot (\mathbf{X} \mathbf{w} - \mathbf{y}) \\ &= \mathbf{X} \mathbf{w} \cdot \mathbf{X} \mathbf{w} - 2\mathbf{X} \mathbf{w} \cdot \mathbf{y} + \mathbf{y} \cdot \mathbf{y} \\ &= (\mathbf{X} \mathbf{w})^T \mathbf{X} \mathbf{w} - 2(\mathbf{X} \mathbf{w})^T \mathbf{y} + \mathbf{y}^T \mathbf{y} \\ &= \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{w}^T \mathbf{X}^T \mathbf{y} + \mathbf{y}^T \mathbf{y}, \end{aligned}$$

and then we'd need to differentiate those terms w.r.t. \mathbf{w} . However, a better way is to use the chain rule. Define f and g such that $J : \mathbb{R}^d \rightarrow \mathbb{R}$ is their composition $J = g \circ f$:

$$\begin{aligned} f : \mathbb{R}^d \rightarrow \mathbb{R}^n & & f(\mathbf{w}) &= \mathbf{X} \mathbf{w} - \mathbf{y} \\ g : \mathbb{R}^n \rightarrow \mathbb{R} & & g(\mathbf{z}) &= \|\mathbf{z}\|^2. \end{aligned}$$

The chain rule says that $\nabla(g \circ f) = (Df)^T \nabla g$, where Df is the derivative of f , i.e. the Jacobian matrix of first partial derivatives¹². We have $Df(\mathbf{w}) = \mathbf{X}$ and $\nabla g(\mathbf{z}) = 2\mathbf{z}$, so

$$\begin{aligned} \nabla J(\mathbf{w}) &= 2\mathbf{X}^T (\mathbf{X} \mathbf{w} - \mathbf{y}) \\ &= 2\mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{X}^T \mathbf{y}. \end{aligned}$$

Penalized Regression

TODO

Logistic Regression

- Two classes.
- The observations y_i are class labels (or probabilities thereof).

¹²The gradient ∇ applies only to scalar-valued functions.

- The model states that the probability of being in class 1 is given by the usual linear model, mapped onto $(0, 1)$ by the logistic function s :

$$\begin{aligned} y_i &\sim \text{Bern}(s(\mathbf{x}_i^T \mathbf{w})), \\ s(z) &= \frac{1}{1 + e^{-z}}. \end{aligned}$$

Note that $s'(z) = \frac{e^{-z}}{(1+e^{-z})^2} = s(z)(1 - s(z))$.

Likelihood

Let $s_i = s(\mathbf{x}_i^T \mathbf{w})$.

$$\begin{aligned} \mathcal{L}(\mathbf{w}) &= \prod_i s_i^{y_i} (1 - s_i)^{(1-y_i)} \\ l(\mathbf{w}) &= \sum_i y_i \log s_i + (1 - y_i) \log(1 - s_i) \\ \nabla l(\mathbf{w}) &= \sum_i \frac{y_i}{s_i} (s_i)(1 - s_i) \mathbf{x}_i + \frac{1 - y_i}{1 - s_i} (-1) (s_i)(1 - s_i) \mathbf{x}_i \\ &= \sum_i \mathbf{x}_i (y_i(1 - s_i) - (1 - y_i)s_i) \\ &= \sum_i \mathbf{x}_i (y_i - s_i) \\ &= \mathbf{X}^T (\mathbf{y} - s(\mathbf{X} \mathbf{w})) \quad (d \times 1) \end{aligned}$$

where $s : \mathbb{R}^n \rightarrow \mathbb{R}^n$ applies s componentwise to the rows.

Optimization problem: Find \mathbf{w} that minimizes the cost function $J(\mathbf{w}) = -l(\mathbf{w})$.

Because the weights \mathbf{w} are tied up inside $s_i = s(\mathbf{x}_i^T \mathbf{w})$ it's not possible to find the minimum \mathbf{w}^* by setting the gradient equal to zero (i.e. by solving a linear system). We can use gradient descent, or Newton's method.

For Newton's method, we need the Hessian of the objective function. This is the $d \times d$ matrix of partial derivatives of the gradient, i.e. \mathbf{X}^T multiplied by the derivative (Jacobian matrix) of $s(\mathbf{X} \mathbf{w})$. Define $\mathbf{f}(\mathbf{w}) = \mathbf{X} \mathbf{w}$ so now $s(\mathbf{X} \mathbf{w}) = (s \circ \mathbf{f})(\mathbf{w})$.

Function	domain \rightarrow range	Jacobian	dim Jacobian
$\mathbf{f}(\mathbf{w}) = \mathbf{X} \mathbf{w}$	$\mathbb{R}^d \rightarrow \mathbb{R}^n$	$D\mathbf{f} = \mathbf{X}$	$n \times d$
$s(\mathbf{z})$	$\mathbb{R}^n \rightarrow \mathbb{R}^n$	$Ds(\mathbf{z}) = \mathbf{S}$	$n \times n$

where \mathbf{S} is a diagonal matrix with $S_{ii} = s(\mathbf{x}_i^T \mathbf{w})(1 - s(\mathbf{x}_i^T \mathbf{w}))$. Now by the chain rule,

$$\begin{aligned} \nabla^2 J(\mathbf{w}) &= \mathbf{X}^T D_s s(\mathbf{X} \mathbf{w}) \\ &= \mathbf{X}^T (D_s s)(D_s \mathbf{f}) \\ &= \mathbf{X}^T \mathbf{S} \mathbf{X}. \end{aligned}$$

1 Change of basis

Suppose person B uses some other basis vectors to describe locations in space. Specifically, in our coordinates, their basis vectors are $\begin{bmatrix} 2 \\ 1 \end{bmatrix}$ and $\begin{bmatrix} -1 \\ 1 \end{bmatrix}$.

When they state a vector, what is it in our coordinates?

If they say $\begin{bmatrix} 2 \\ 1 \end{bmatrix}$, what is that in our coordinates?

Well, if they say $\begin{bmatrix} 2 \\ 1 \end{bmatrix}$, that's $\begin{bmatrix} 2 \\ 1 \end{bmatrix}$ in our coordinates. And if they say $\begin{bmatrix} -1 \\ 1 \end{bmatrix}$, that's $\begin{bmatrix} -1 \\ 1 \end{bmatrix}$ in our coordinates. So the matrix containing their *basis vectors expressed using our coordinate system* transforms a point expressed in their coordinate system into one expressed in ours. That last sentence is critical, so hopefully it makes sense! So, the answer is

$$\mathbf{2-111} \begin{bmatrix} -1 \\ 2 \end{bmatrix} = \begin{bmatrix} -4 \\ 1 \end{bmatrix}.$$

When we state a vector, what is it in their coordinates?

We give the vector $\begin{bmatrix} 2 \\ 1 \end{bmatrix}$. What is that in their coordinate system? By definition, the answer is the weights that scales their basis vectors to hit $\begin{bmatrix} 2 \\ 1 \end{bmatrix}$. So, the solution to

$$\mathbf{2-111} \begin{bmatrix} a \\ b \end{bmatrix} = \begin{bmatrix} 3 \\ 2 \end{bmatrix}.$$

Computationally, we can see that we can get the solution by multiplying both sides by the inverse:

$$\begin{bmatrix} a \\ b \end{bmatrix} = \mathbf{2-111}^{-1} \begin{bmatrix} 3 \\ 2 \end{bmatrix}.$$

Conceptually, we have

$$\mathbf{2-111} = \begin{bmatrix} \text{matrix converting their} \\ \text{representation to ours} \end{bmatrix}$$

where “their representation” means the vector expressed using their coordinate system. So the role played by the inverse is

$$\begin{bmatrix} a \\ b \end{bmatrix} = \begin{bmatrix} \text{matrix converting our} \\ \text{representation to theirs} \end{bmatrix} \begin{bmatrix} 3 \\ 2 \end{bmatrix}.$$

When we state a transformation, what is it in their coordinates?

We state a 90 anticlockwise rotation of 2D space:

$$\mathbf{0-110}$$

what is that transformation in their coordinates? The answer is

$$\begin{bmatrix} \text{matrix converting our} \\ \text{representation to theirs} \end{bmatrix} \mathbf{0-110} \begin{bmatrix} \text{matrix converting their} \\ \text{representation to ours} \end{bmatrix}$$

since the composition of those three transformations defines a single transformation that takes in a vector expressed in their coordinate system, converts it to our coordinate system, transforms it as requested, and then converts back to theirs.

2 Symmetric matrices

Spectral theorem for symmetric matrices

Symmetric $n \times n$ matrix A (real).

$$A^{-1} = A^T$$

n orthogonal eigenvectors with real eigenvalues.

Orthonormal matrix U containing normalized eigenvectors.

$$A = U M U^{-1} = U M U^T$$

(Eigenvalues are uniquely determined by matrix. Eigenvalues can be repeated, in which case any linear combination of their eigenvectors is also an eigenvector.)

Linear and quadratic approximations to a function¹³

We construct first- and second-order approximations to a differentiable function $f : \mathbb{R}^2 \rightarrow \mathbb{R}$. The approximation is made at some point $(x_0, y_0) = \mathbf{x}_0 \in \mathbb{R}^2$; we demand that the value of the approximation, and the first and second derivatives, match those of f exactly at that point.

Linear approximation to a function $f(x, y)$ near (x_0, y_0) :

$$\begin{aligned} L(x, y) &= f(\mathbf{x}_0, y_0) + (x - x_0) f_x(\mathbf{x}_0, y_0) + (y - y_0) f_y(\mathbf{x}_0, y_0) \\ &= f(\mathbf{x}) + (\mathbf{x} - \mathbf{x}_0) \cdot \nabla f(\mathbf{x}_0) \end{aligned}$$

Note that, at (x_0, y_0) , the first partial derivatives of L are equal to those of f , as they must be. (In fact, we could say that the coefficients are determined by this requirement; see the quadratic case below. But the linear case is obvious without “deriving” the coefficients.)

Quadratic approximation to a function $f(x, y)$ near (x_0, y_0) :

First note that the “quadratic form” $a x^2 + 2bxy + c y^2$ can be written as

$$\mathbf{x}^T A \mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}^T \text{abdc} \begin{bmatrix} x \\ y \end{bmatrix}$$

¹³ *khanacademy - Grant Sanderson - second partial derivative test*

This is a scalar. In general, a quadratic form for symmetric matrix A is

$$\mathbf{x}^T A \mathbf{y} = \sum_{j,k} A_{j,k} x_j y_k.$$

The j -th component of the gradient of $q(\mathbf{x}) = \mathbf{x}^T A \mathbf{x}$ is $\frac{\partial q}{\partial x_j} = 2 \sum_k A_{j,k} x_k$, so

$$\nabla \mathbf{x}^T A \mathbf{x} = 2A \mathbf{x}.$$

$$\begin{aligned} Q(x, y) &= f(\mathbf{x}_0) + (x - x_0) f_x(\mathbf{x}_0) + (y - y_0) f_y(\mathbf{x}_0) + \\ &\quad \frac{1}{2} f_{xx}(\mathbf{x}_0) (x - x_0)^2 + f_{xy}(\mathbf{x}_0) (x - x_0) (y - y_0) + \frac{1}{2} f_{yy}(\mathbf{x}_0) (y - y_0)^2 \\ &= f(\mathbf{x}_0) + (\mathbf{x} - \mathbf{x}_0) \cdot \nabla f(\mathbf{x}_0) + \frac{1}{2} (\mathbf{x} - \mathbf{x}_0)^T \nabla^2 f(\mathbf{x}_0) (\mathbf{x} - \mathbf{x}_0), \end{aligned}$$

where $\nabla^2 f(\mathbf{x}_0)$ is the Hessian matrix $f_{xx} f_{xy} f_{yx} f_{yy}$ evaluated at \mathbf{x}_0 .

Second partial derivative test and positive definiteness of Hessian

The second partial derivative test for a function of two variables states that we examine the determinant of the Hessian evaluated at the critical point:

$$D = \det \nabla^2 f(\mathbf{x}_0) = f_{xx}(\mathbf{x}_0) f_{yy}(\mathbf{x}_0) - f_{xy}(\mathbf{x}_0)^2.$$

Notice that $D \geq 0$ implies that the sign of f_{xx} and f_{yy} agree (because we're subtracting the square of the mixed partial f_{xy} , i.e. a positive number).

D	roots	f_{xx}	Hessian
+	no real roots	+	minimum
+	no real roots	-	maximum
0	one real root	+	positive semidefinite
0	one real root	-	negative semidefinite
-	two real roots	n/a	saddle point

Explanation

At a critical point \mathbf{x}_0 , the gradient is zero and the quadratic approximation is therefore

$$Q(x, y) = f(\mathbf{x}_0) + \frac{1}{2} (\mathbf{x} - \mathbf{x}_0)^T \nabla^2 f(\mathbf{x}_0) (\mathbf{x} - \mathbf{x}_0).$$

So if this is a minimum (concave-up paraboid) then this quadratic form is positive for all $\mathbf{x} \neq \mathbf{x}_0$ (and if it's a maximum then it's negative for all $\mathbf{x} \neq \mathbf{x}_0$).

Basically the argument is that, instead of analyzing the function f itself, we analyze its quadratic approximation at the critical point. So the question comes down to: how do we determine whether a quadratic form is always positive, always negative, or takes positive and negative values?